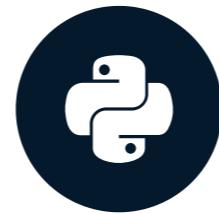


Tensors, layers and autoencoders

INTRODUCTION TO DEEP LEARNING WITH KERAS



Miguel Esteban

Data Scientist & Founder

Accessing Keras layers

```
# Accessing the first layer of a Keras model  
first_layer = model.layers[0]  
# Printing the layer, and its input, output and weights  
print(first_layer.input)  
print(first_layer.output)  
print(first_layer.weights)
```

```
<tf.Tensor 'dense_1_input:0' shape=(?, 3) dtype=float32>
```

```
<tf.Tensor 'dense_1/Relu:0' shape=(?, 2) dtype=float32>
```

```
[<tf.Variable 'dense_1/kernel:0' shape=(3, 2) dtype=float32_ref>,  
<tf.Variable 'dense_1/bias:0' shape=(2,) dtype=float32_ref>]
```

What are tensors?

```
# Defining a rank 2 tensor (2 dimensions)
T2 = [[1,2,3],
      [4,5,6],
      [7,8,9]]

# Defining a rank 3 tensor (3 dimensions)
T3 = [[1,2,3],
      [4,5,6],
      [7,8,9],
      [10,11,12],
      [13,14,15],
      [16,17,18],
      [19,20,21],
      [22,23,24],
      [25,26,27]]
```

```
# Import Keras backend
import keras.backend as K

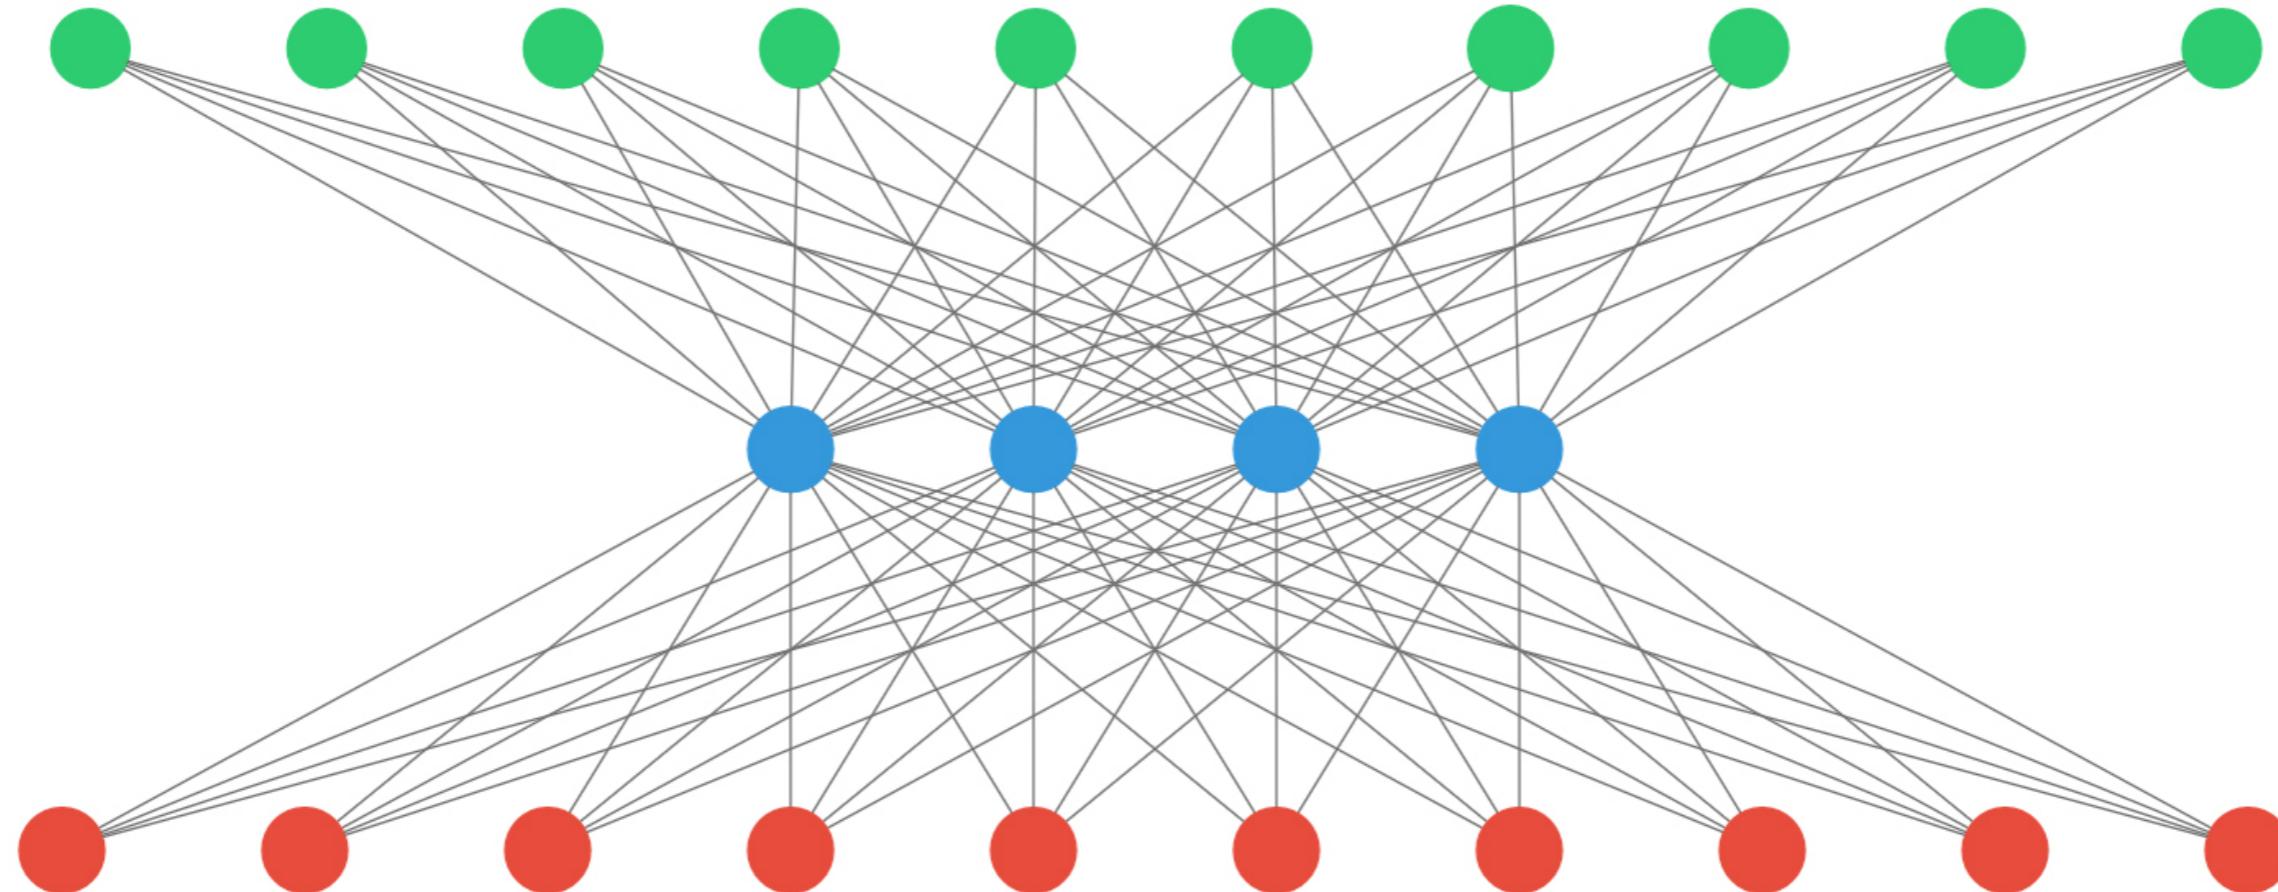
# Get the input and output tensors of a model layer
inp = model.layers[0].input
out = model.layers[0].output
# Function that maps layer inputs to outputs
inp_to_out = K.function([inp], [out])
# We pass an input and get the output we'd get in that first layer
print(inp_to_out([X_train]))
```

```
# Outputs of the first layer per sample in X_train
[array([[0.7, 0], ..., [0.1, 0.3]])]
```

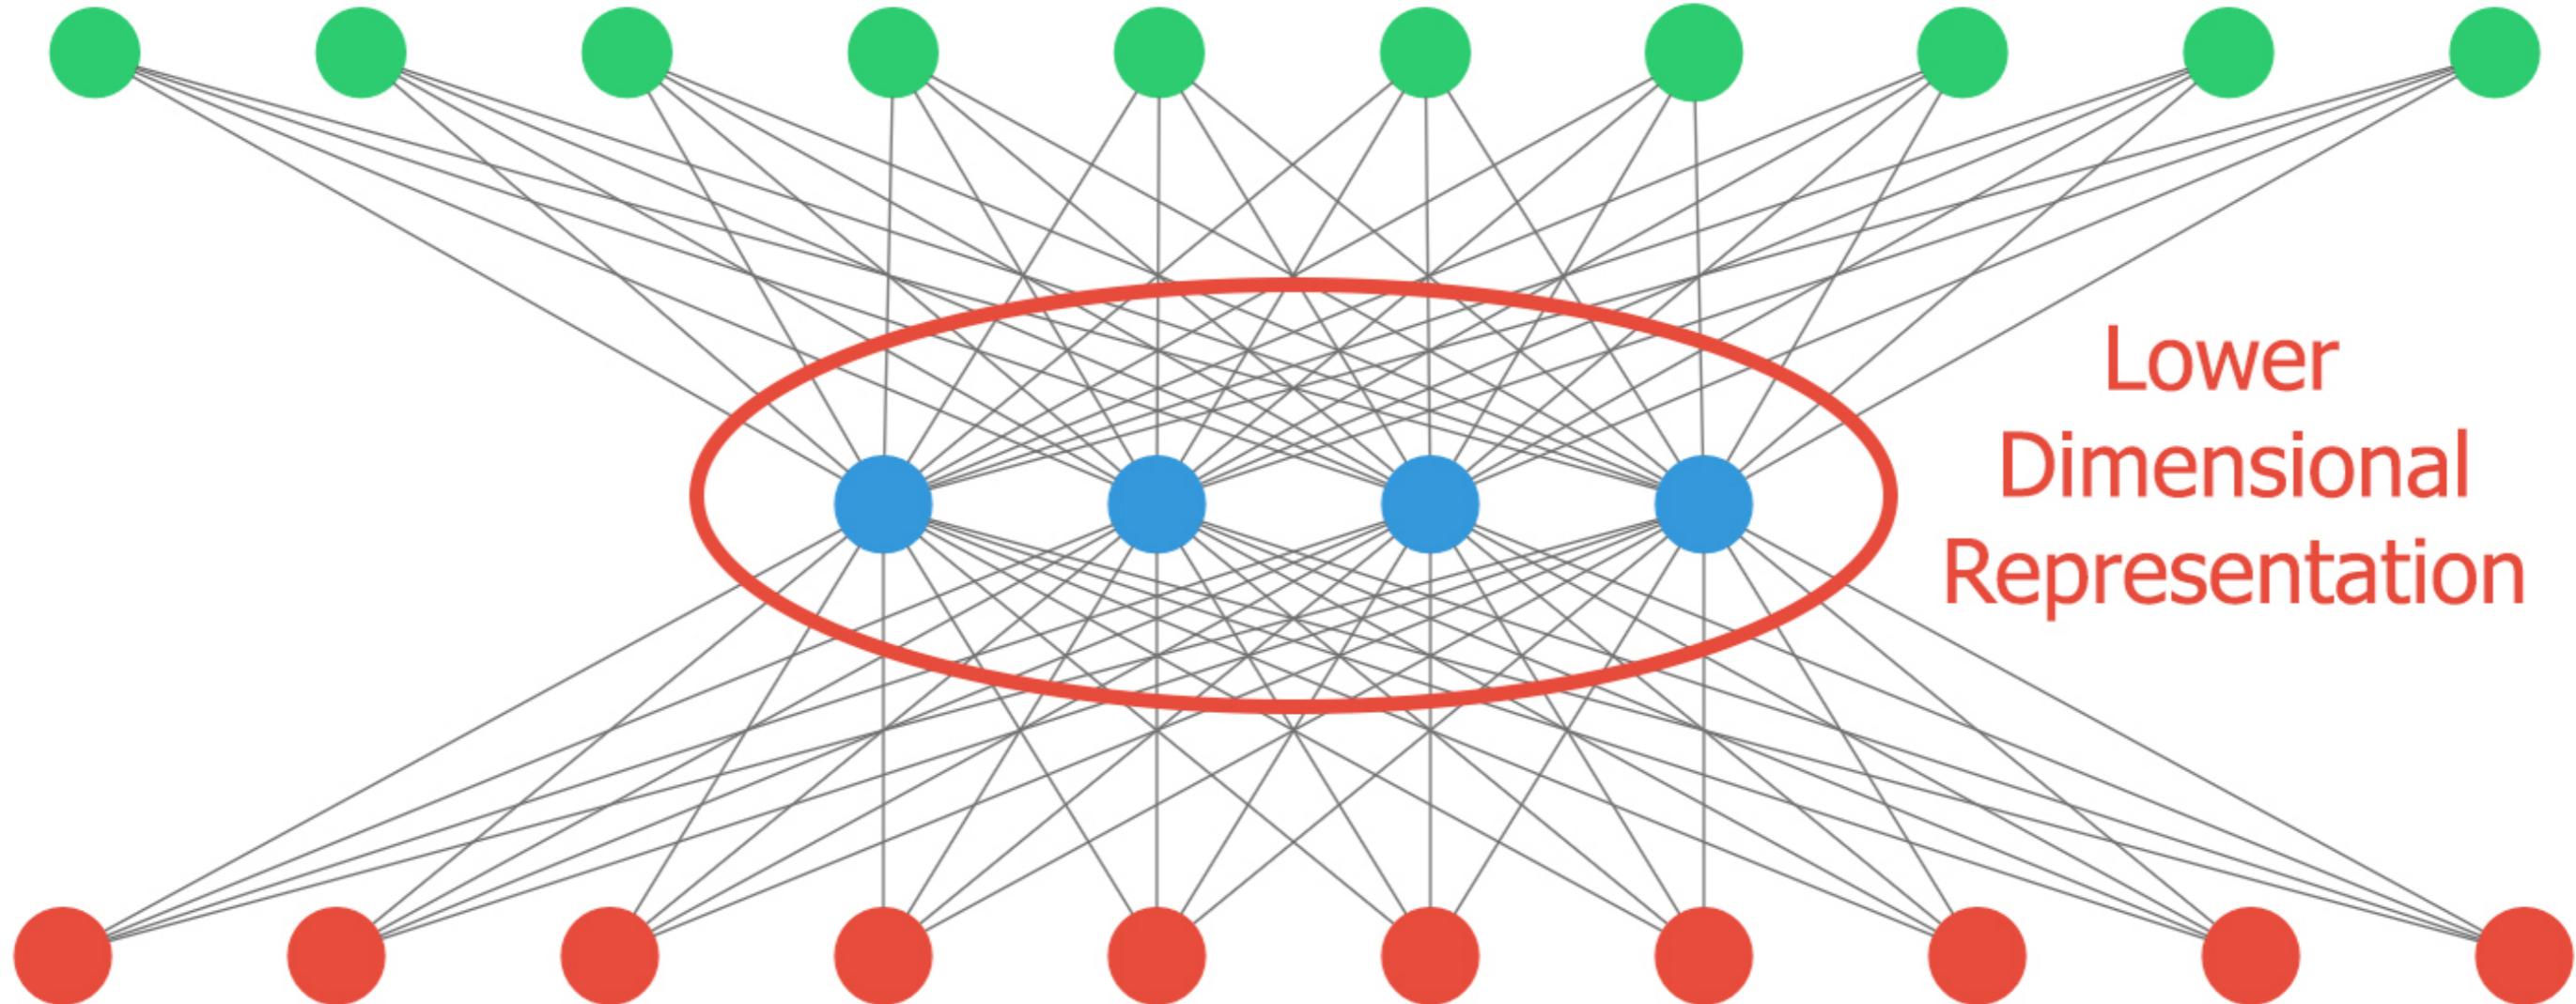
it's time to introduce a new architecture...

Autoencoders!

INPUTS == OUTPUTS



Autoencoders!



Autoencoder use cases

- Dimensionality reduction:
 - Smaller dimensional space representation of our inputs.
- De-noising data:
 - If trained with clean data, irrelevant noise will be filtered out during reconstruction.
- Anomaly detection:
 - A poor reconstruction will result when the model is fed with unseen inputs.
- ...



Building a simple autoencoder

```
# Instantiate a sequential model
autoencoder = Sequential()
# Add a hidden layer of 4 neurons and an input layer of 100
autoencoder.add(Dense(4, input_shape=(100,), activation='relu'))
# Add an output layer of 100 neurons
autoencoder.add(Dense(100, activation='sigmoid'))
# Compile your model with the appropriate loss
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

Breaking it into an encoder

```
# Building a separate model to encode inputs  
encoder = Sequential()  
encoder.add(autoencoder.layers[0])  
# Predicting returns the four hidden layer neuron outputs  
encoder.predict(X_test)
```

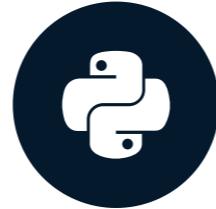
```
# Four numbers for each observation in X_test  
array([10.0234375, 5.833543, 18.90444, 9.20348],...)
```

Let's experiment!

INTRODUCTION TO DEEP LEARNING WITH KERAS

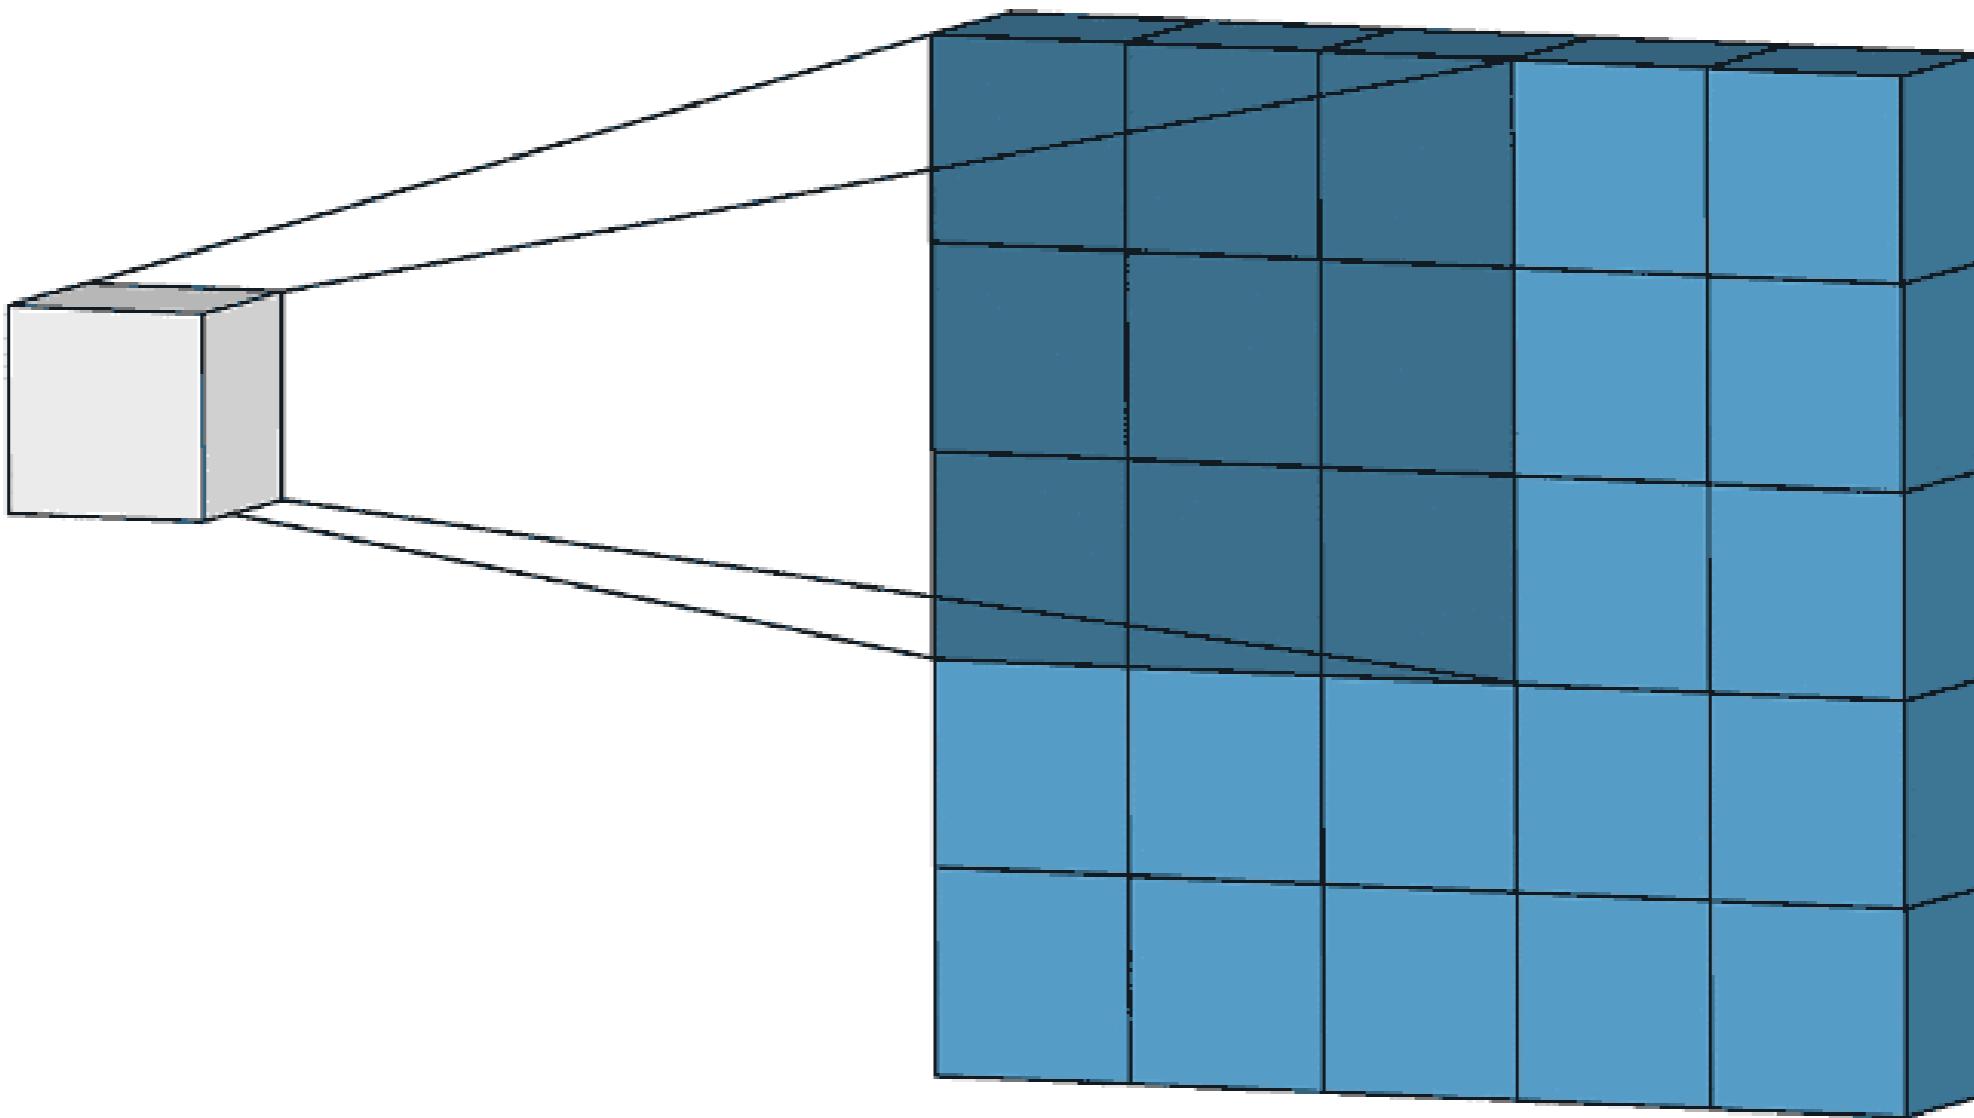
Intro to CNNs

INTRODUCTION TO DEEP LEARNING WITH KERAS

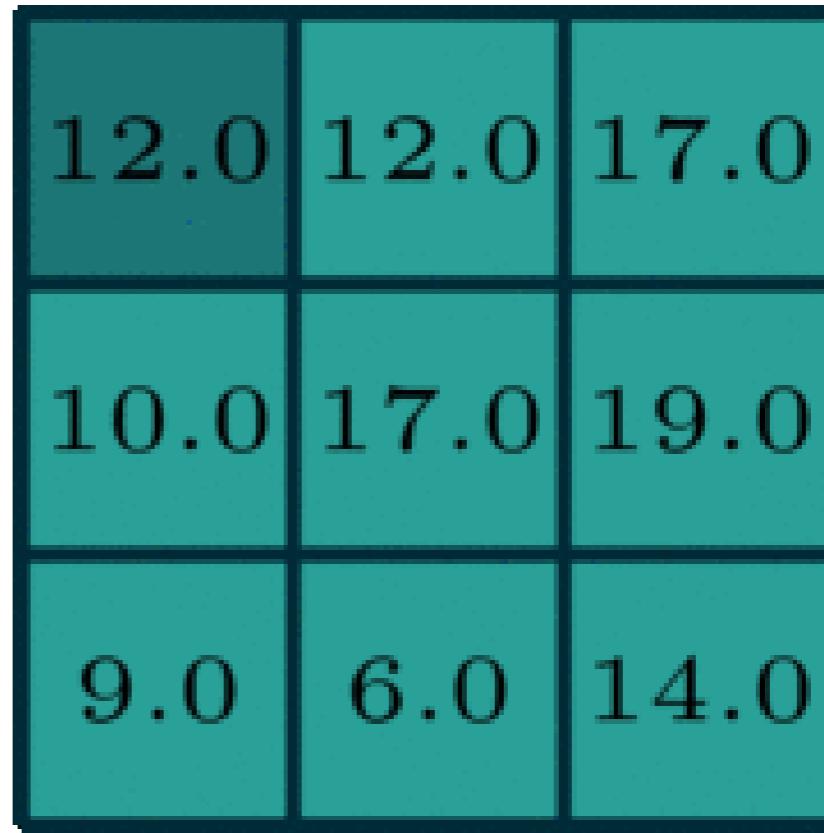


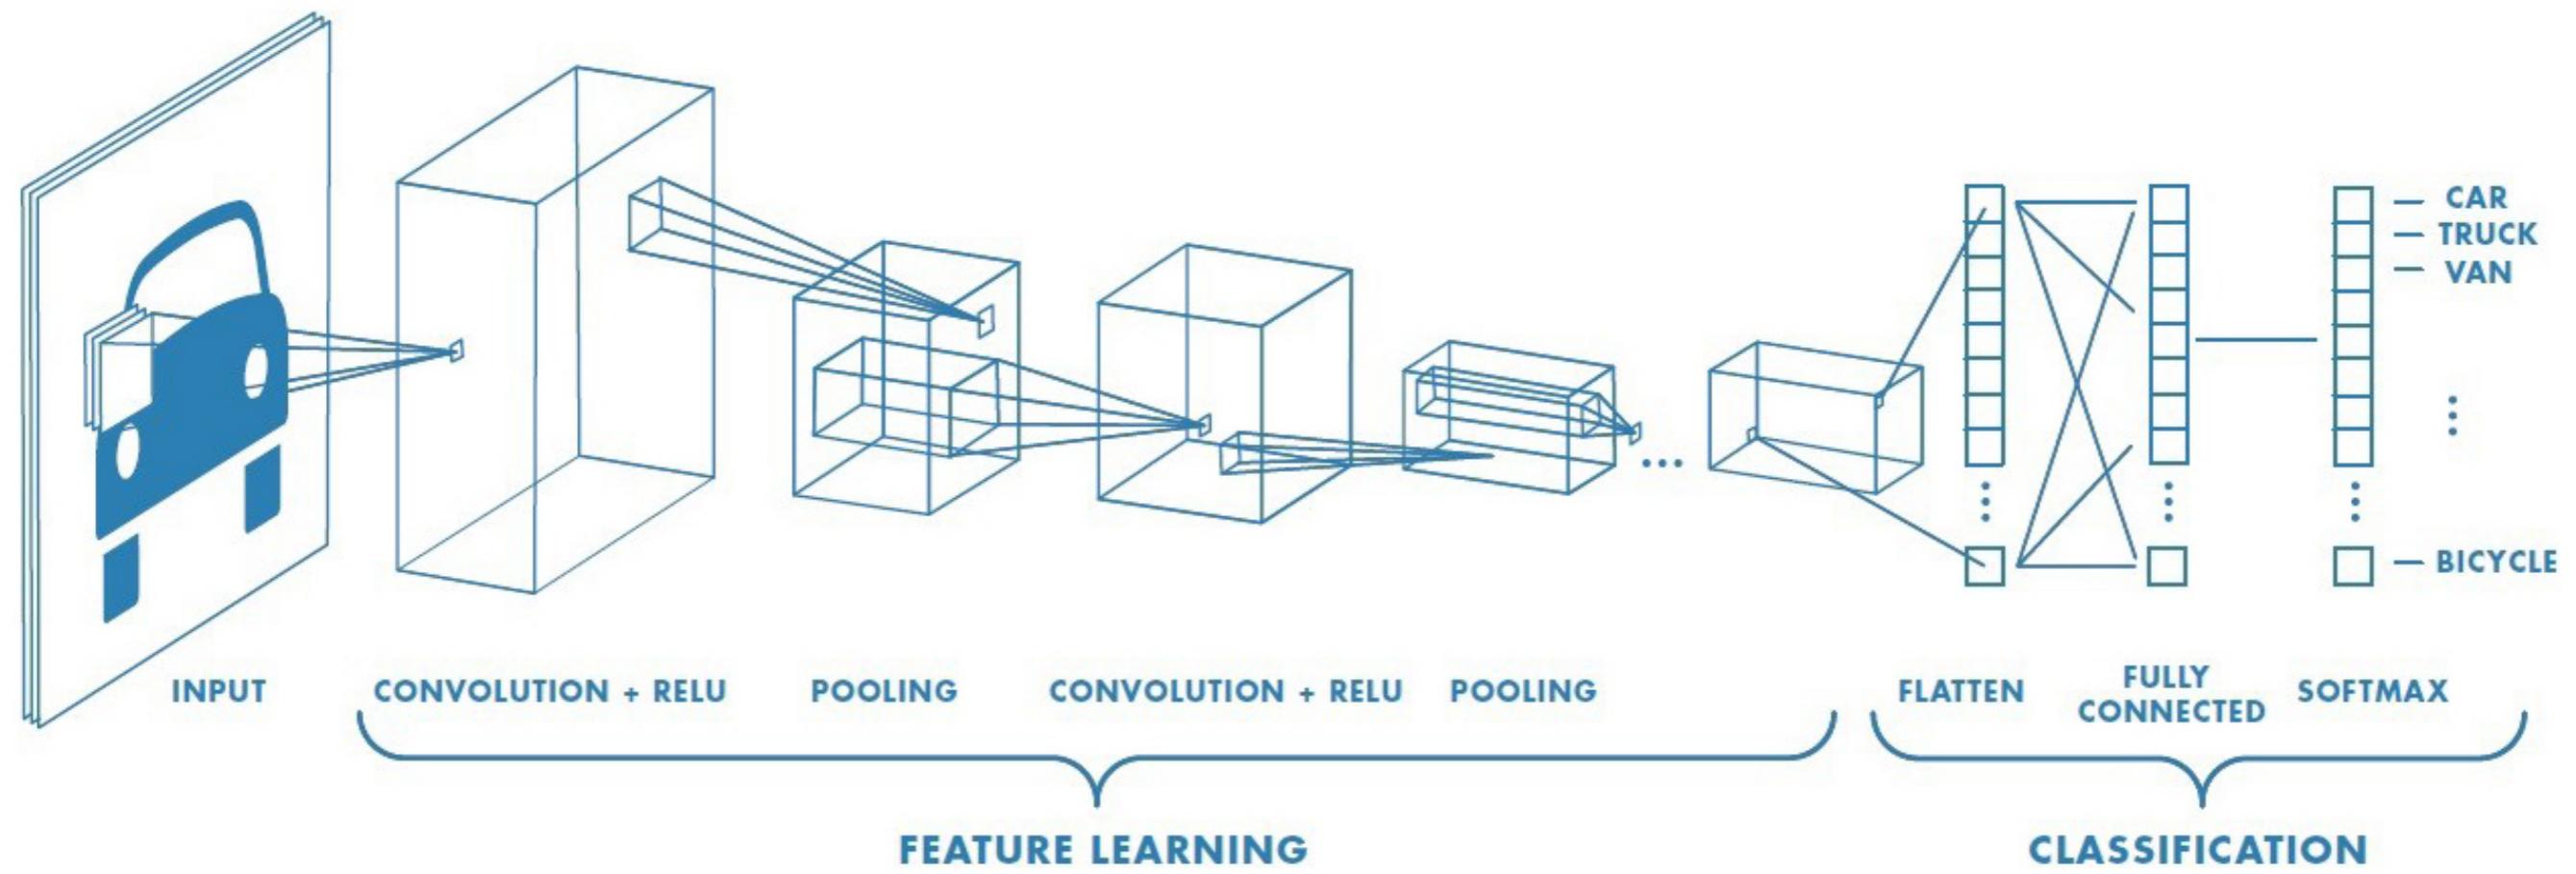
Miguel Esteban

Data Scientist & Founder



3 0	3 1	2 2	1	0
0 2	0 2	1 0	3	1
3 0	1 1	2 2	2	3
2	0	0	2	2
2	0	0	0	1

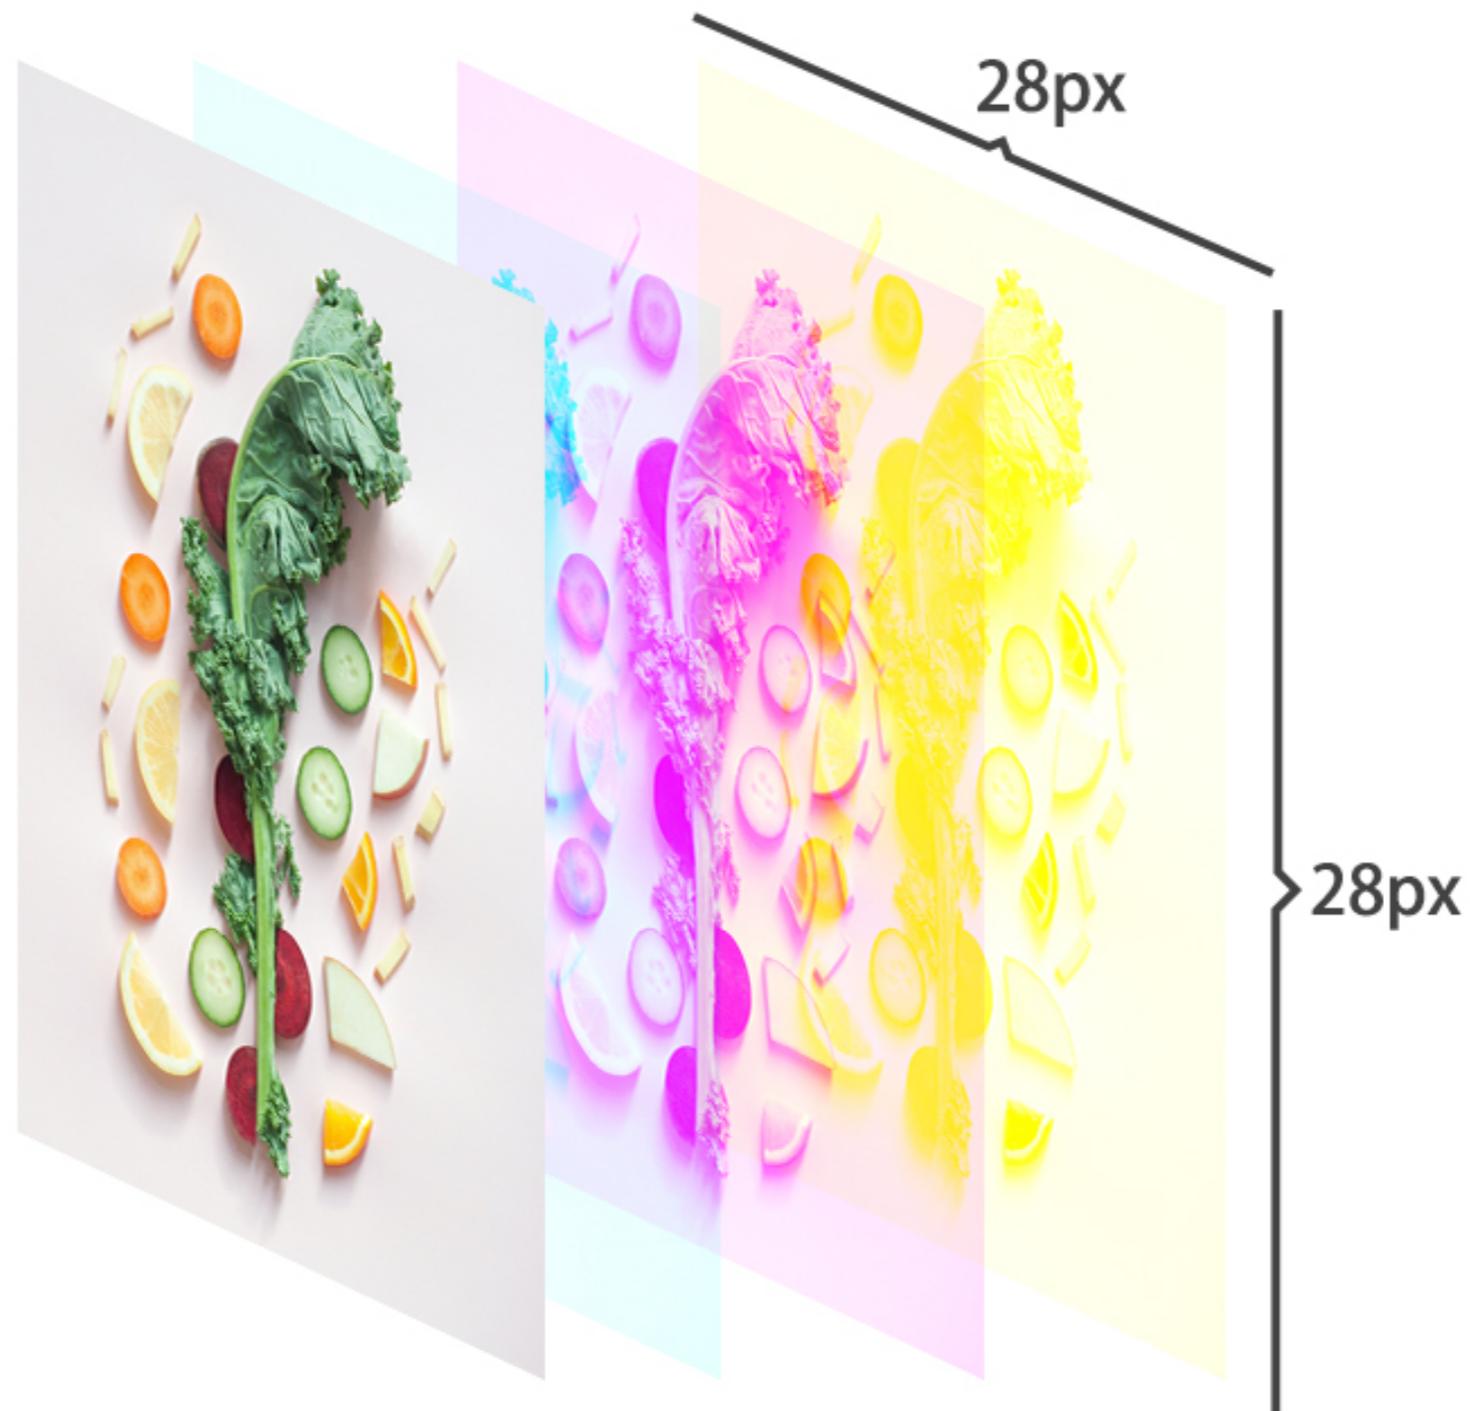




`input_shape`
(WIDTH,HEIGHT,CHANNELS)

`input_shape = (28,28,3)`

Color Images have 3 channels
RGB (Red Green Blue)



```
# Import Conv2D layer and Flatten from keras layers
from keras.layers import Dense, Conv2D, Flatten
# Instantiate your model as usual
model = Sequential()

# Add a convolutional layer with 32 filters of size 3x3
model.add(Conv2D(filters=32,
                  kernel_size=3,
                  input_shape=(28, 28, 1),
                  activation='relu'))

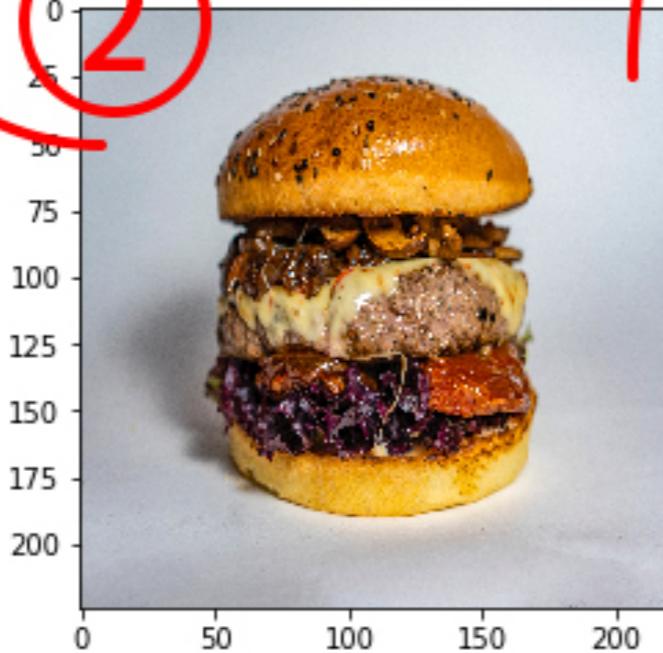
# Add another convolutional layer
model.add(Conv2D(8, kernel_size=3, activation='relu'))

# Flatten the output of the previous layer
model.add(Flatten())
# End this multiclass model with 3 outputs and softmax
model.add(Dense(3, activation='softmax'))
```

1



2



3

ResNet50
(50 Layers)

4

Predicted
("cheeseburger", 0.99),
("bagel", 0.01),
("bakery", 0.00)

Pre-processing images for ResNet50

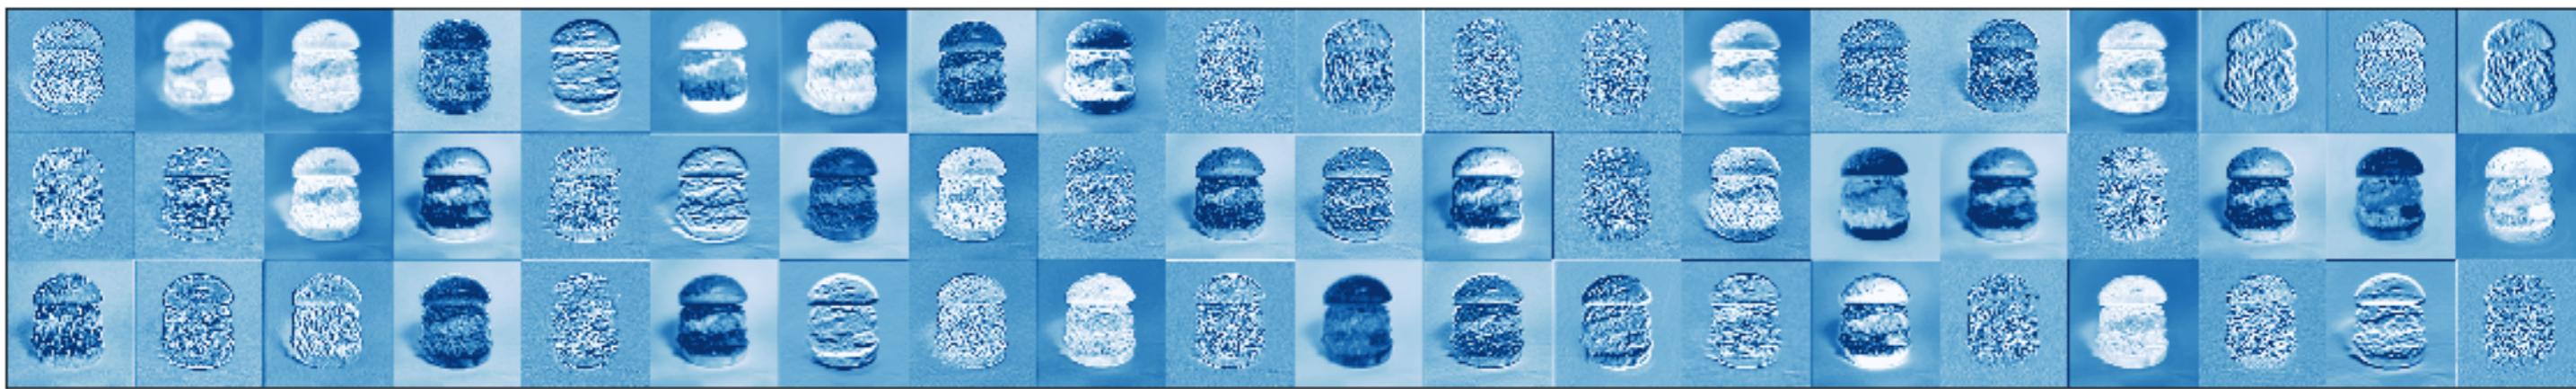
```
# Import image from keras preprocessing
from keras.preprocessing import image
# Import preprocess_input from keras applications resnet50
from keras.applications.resnet50 import preprocess_input
# Load the image with the right target size for your model
img = image.load_img(img_path, target_size=(224, 224))
# Turn it into an array
img = image.img_to_array(img)
# Expand the dimensions so that it's understood by our network:
# img.shape turns from (224, 224, 3) into (1, 224, 224, 3)
img = np.expand_dims(img, axis=0)
# Pre-process the img in the same way training images were
img = preprocess_input(img)
```

Using the ResNet50 model in Keras

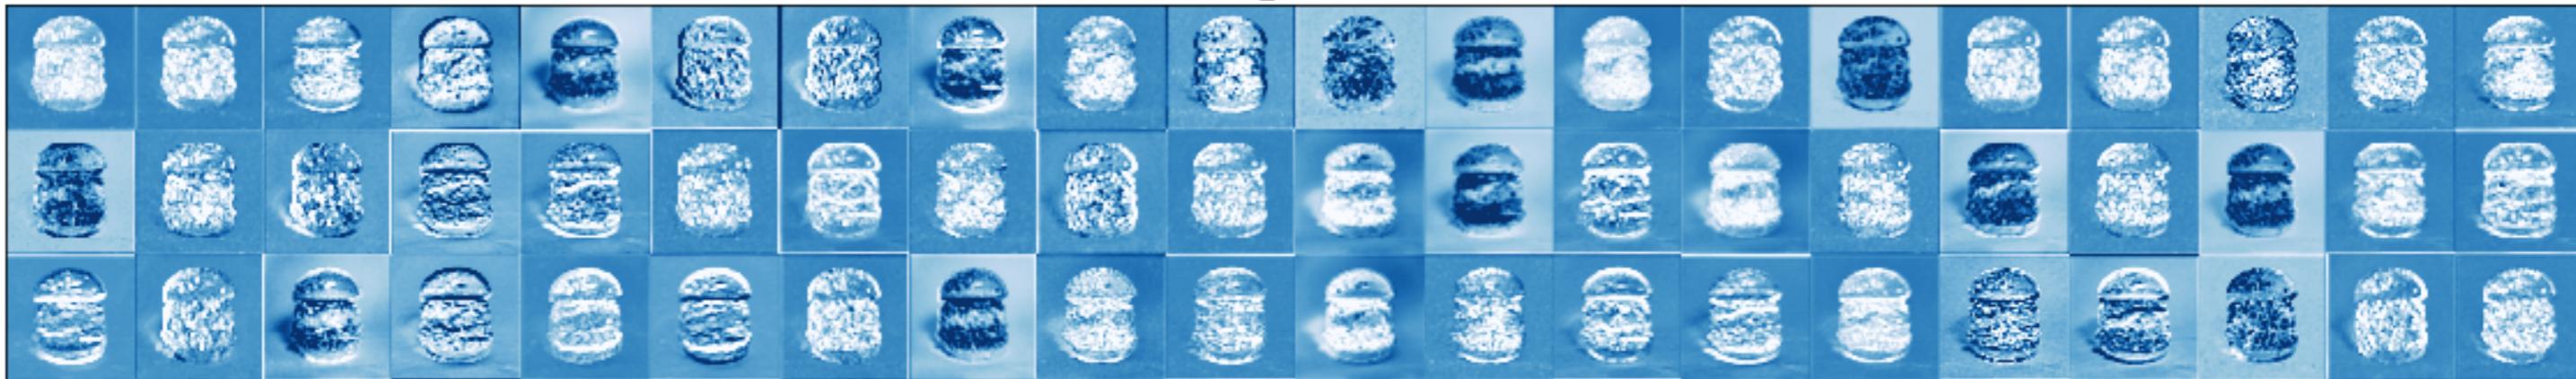
```
# Import ResNet50 and decode_predictions from keras.applications.resnet50
from keras.applications.resnet50 import ResNet50, decode_predictions
# Instantiate a ResNet50 model with imagenet weights
model = ResNet50(weights='imagenet')
# Predict with ResNet50 on our img
preds = model.predict(img)
# Decode predictions and print it
print('Predicted:', decode_predictions(preds, top=1)[0])
```

Predicted: [('n07697313', 'cheeseburger', 0.9868016)]

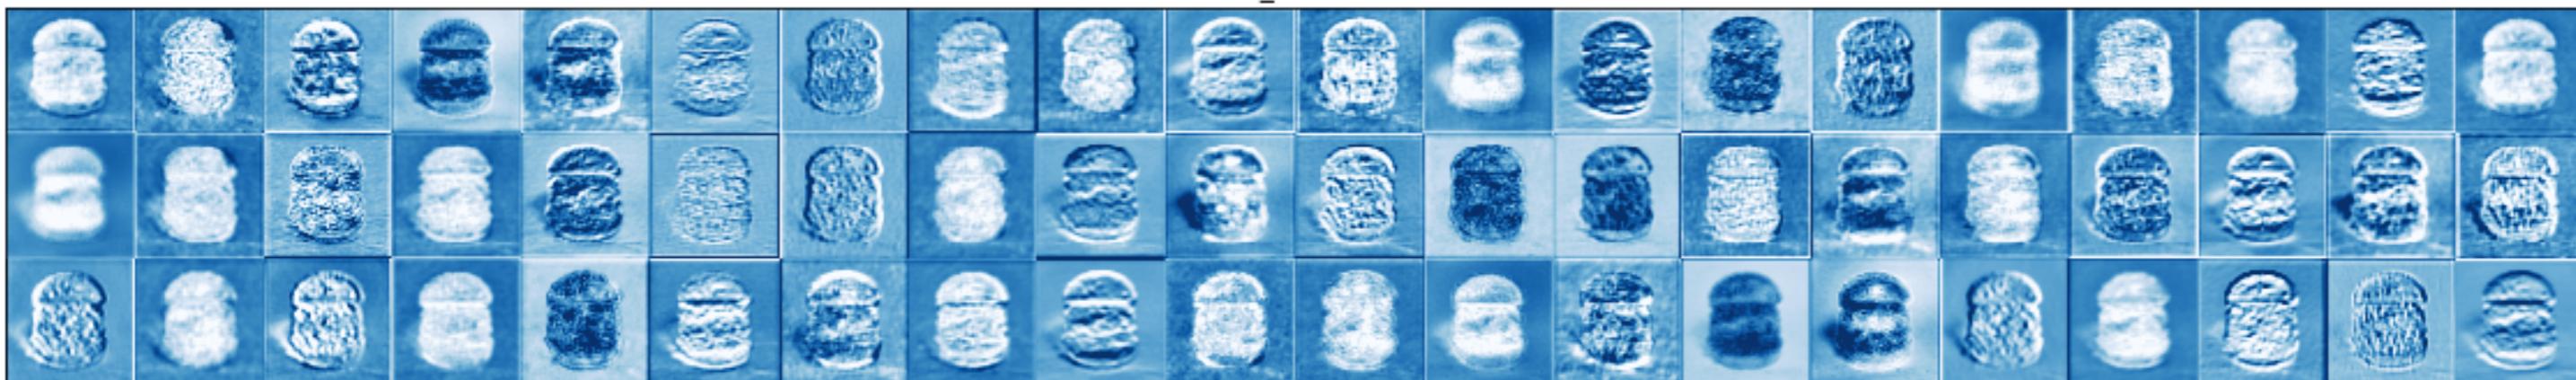
conv1



res2a_branch2a



res2a_branch2b



Let's experiment!

INTRODUCTION TO DEEP LEARNING WITH KERAS

Intro to LSTMs

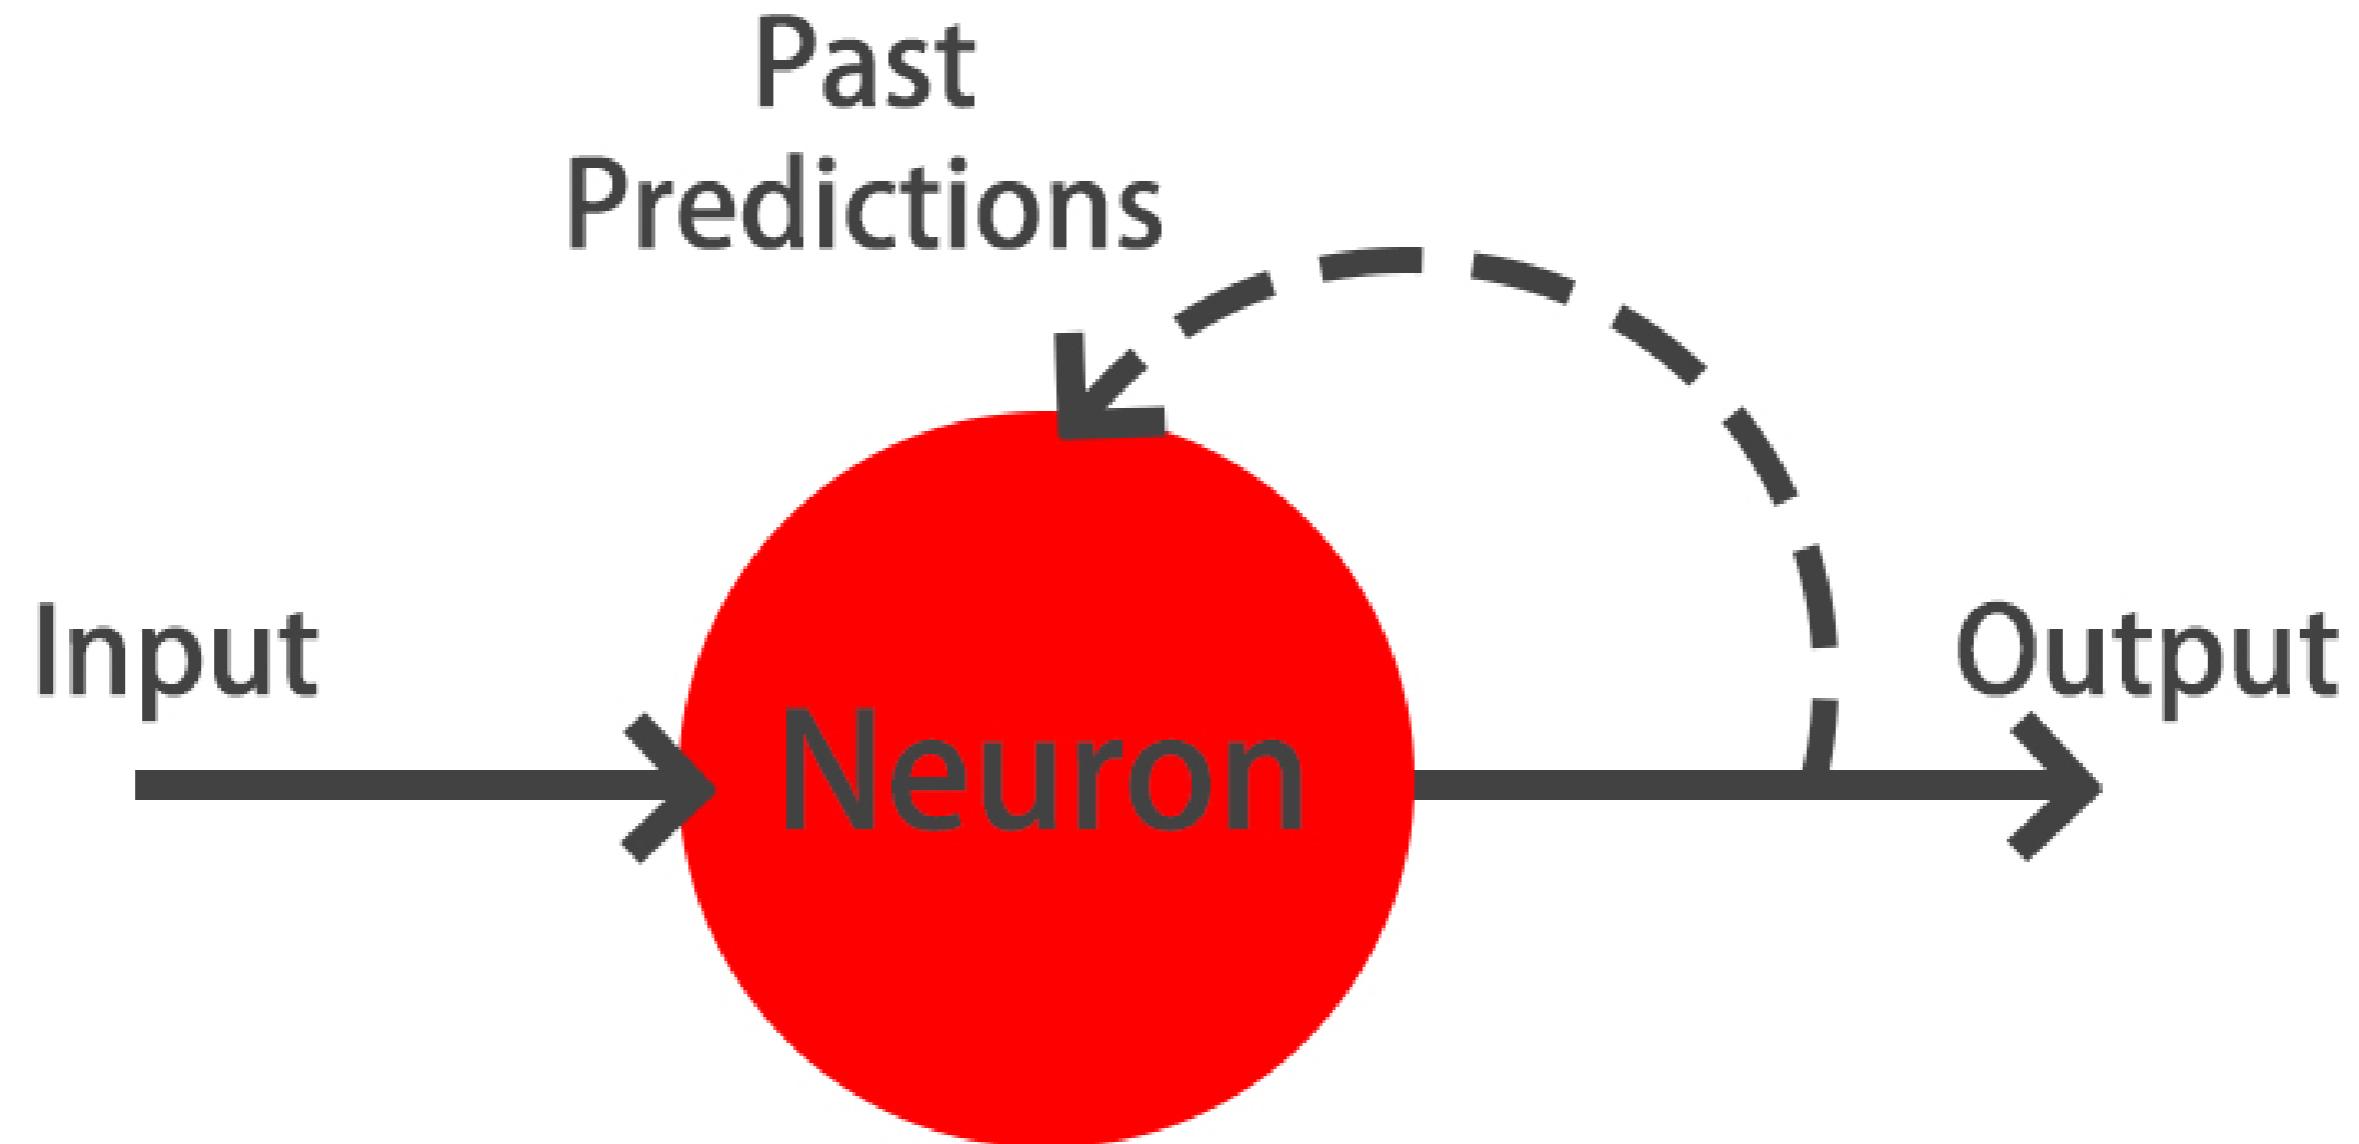
INTRODUCTION TO DEEP LEARNING WITH KERAS

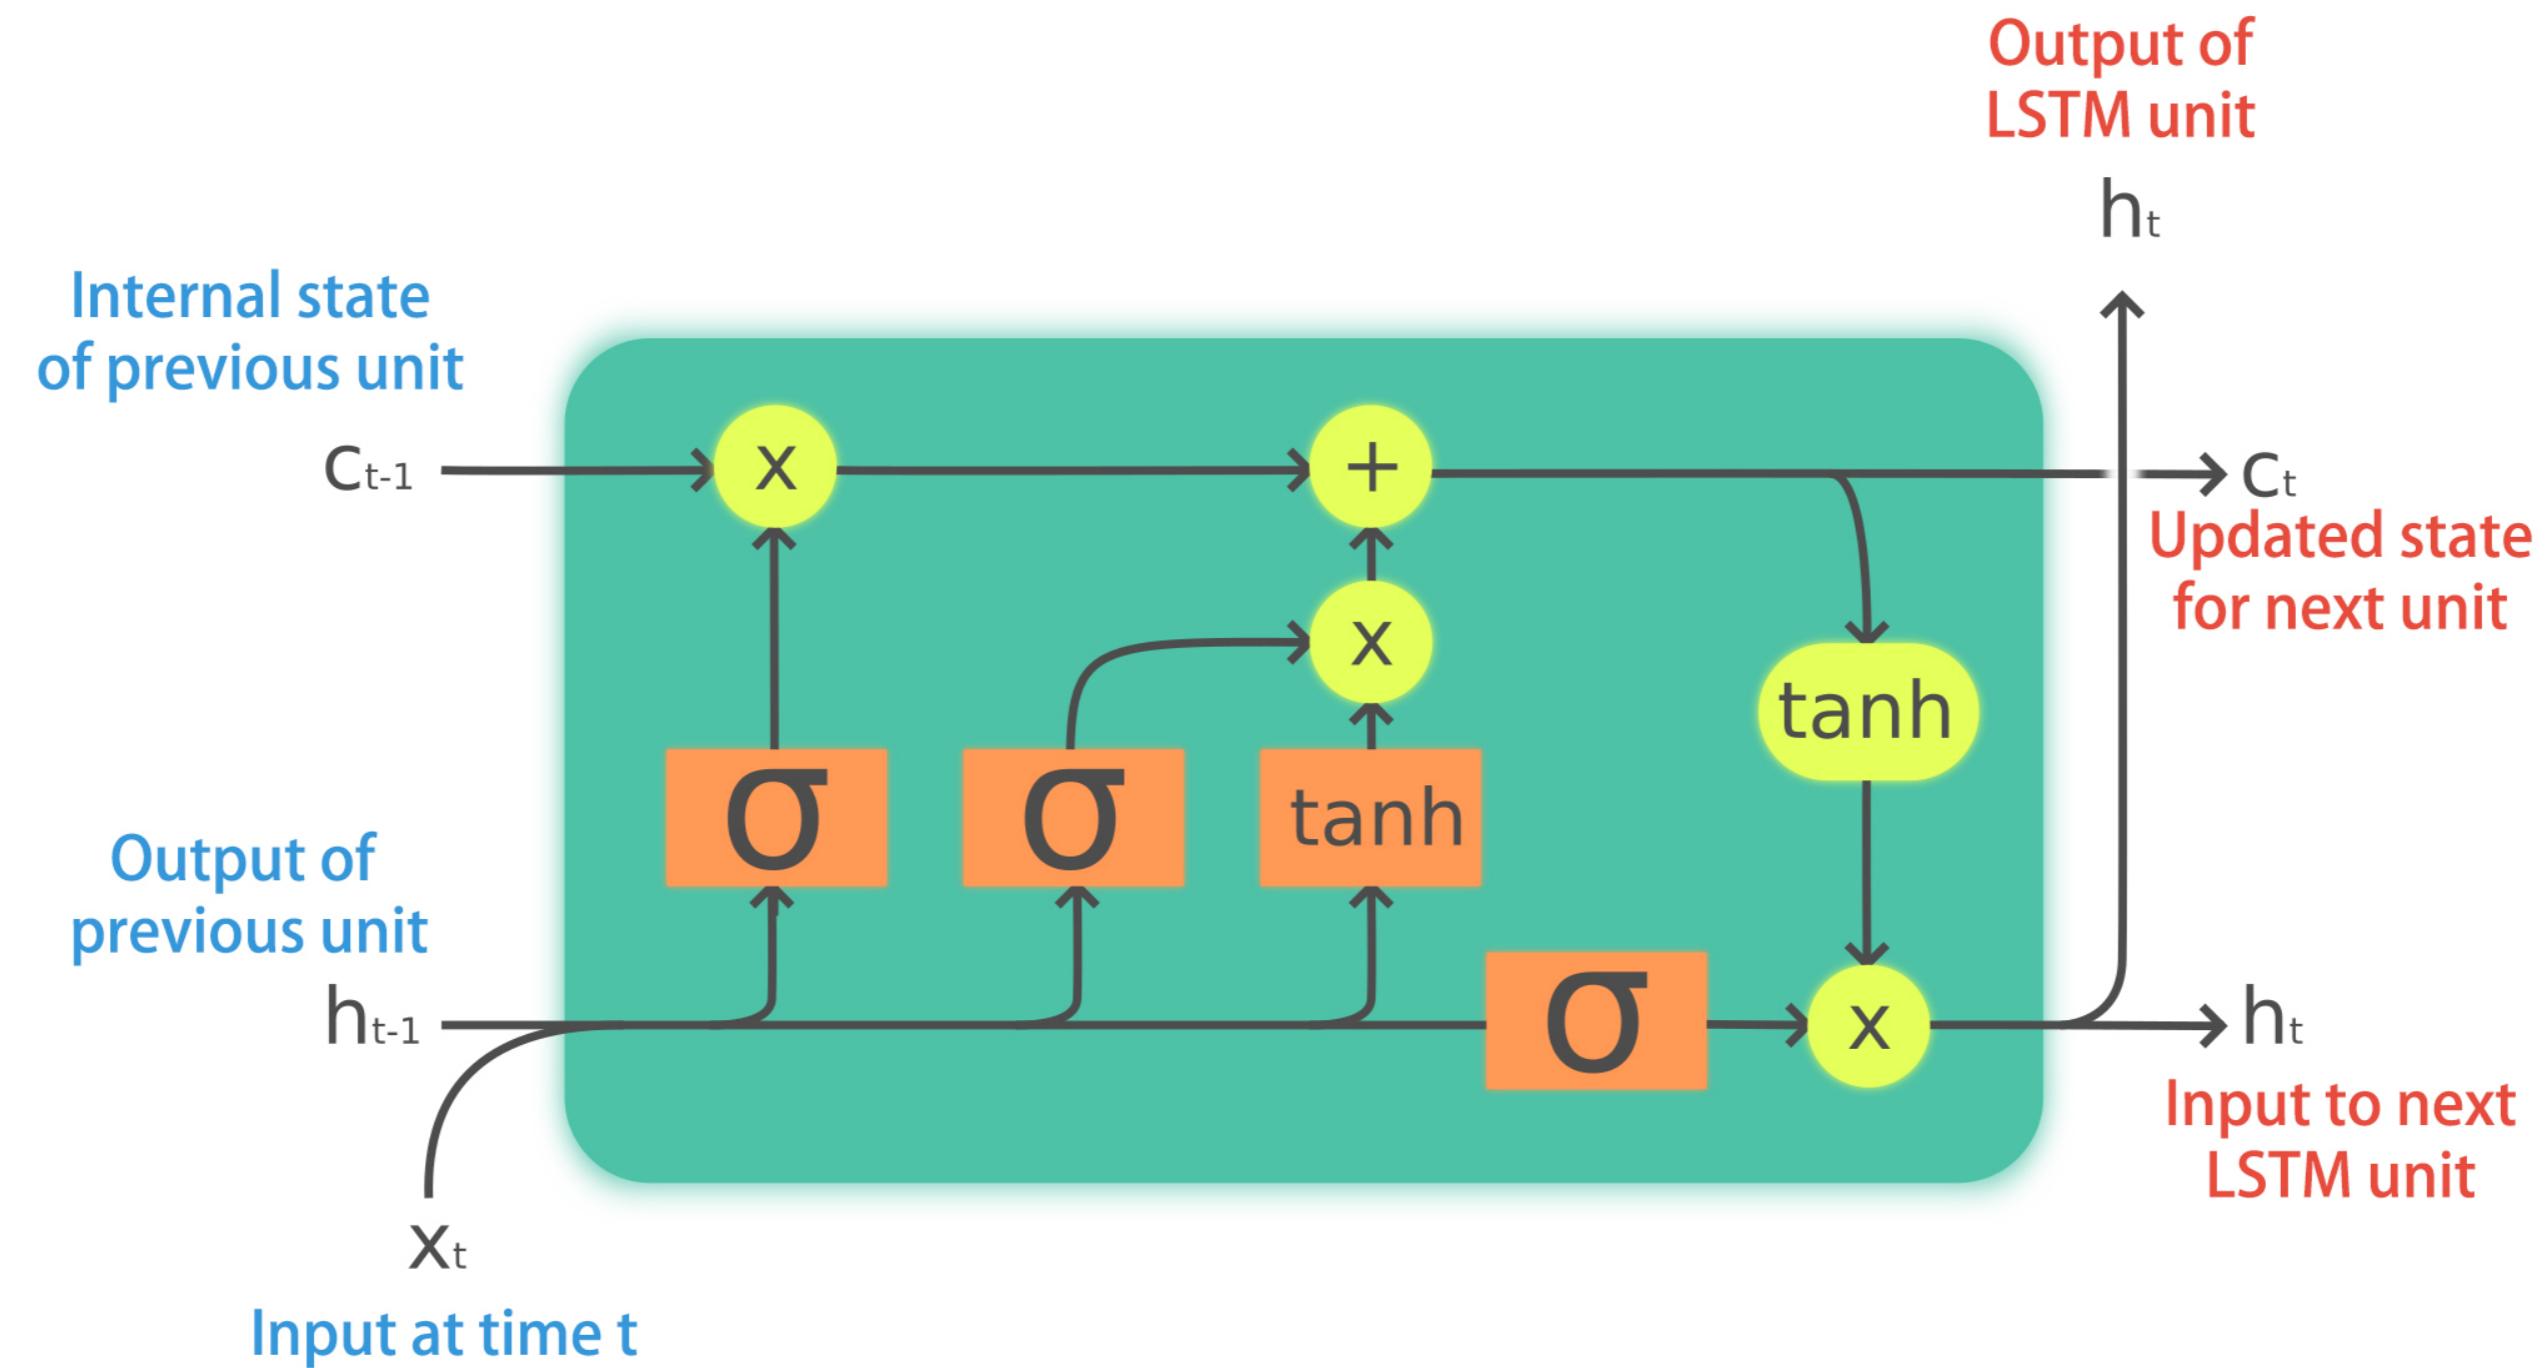


Miguel Esteban

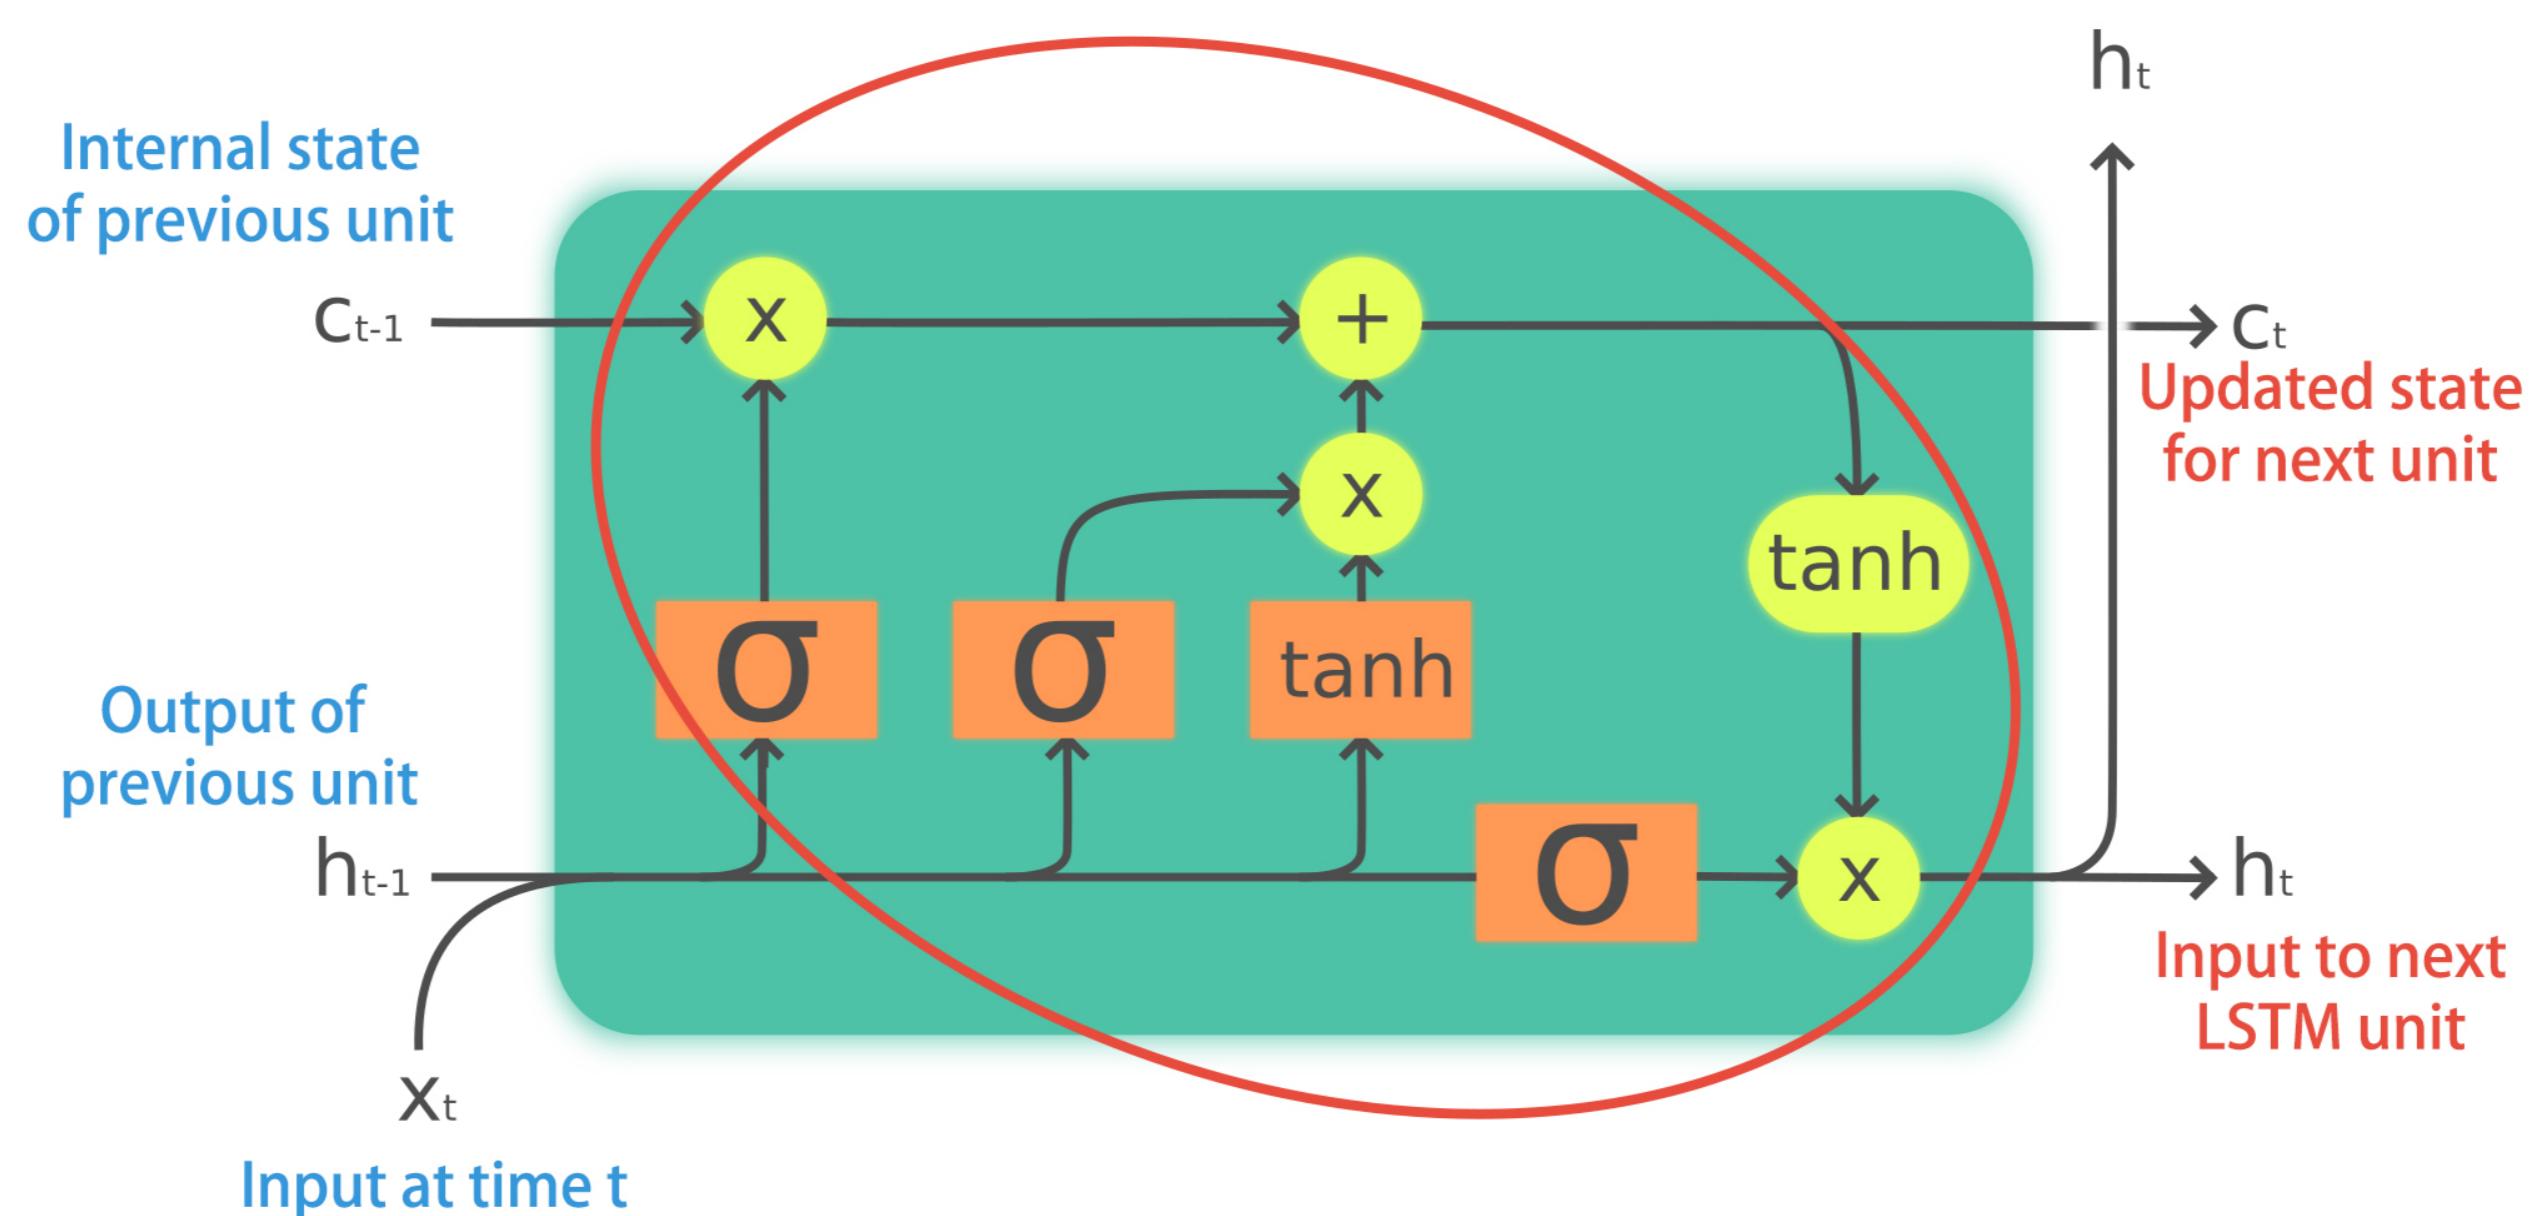
Data Scientist & Founder

What are RNNs?





They learn what to ignore, what to keep
and to select the most important pieces
of past information.

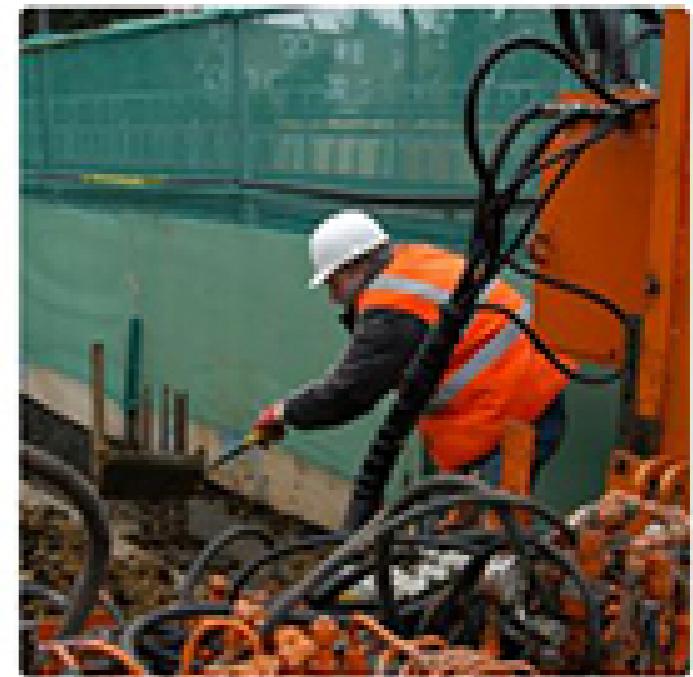


When to use LSTMs?

- Image captioning
- Speech to text
- Text translation
- Document summarization
- Text generation
- Musical composition
- ...

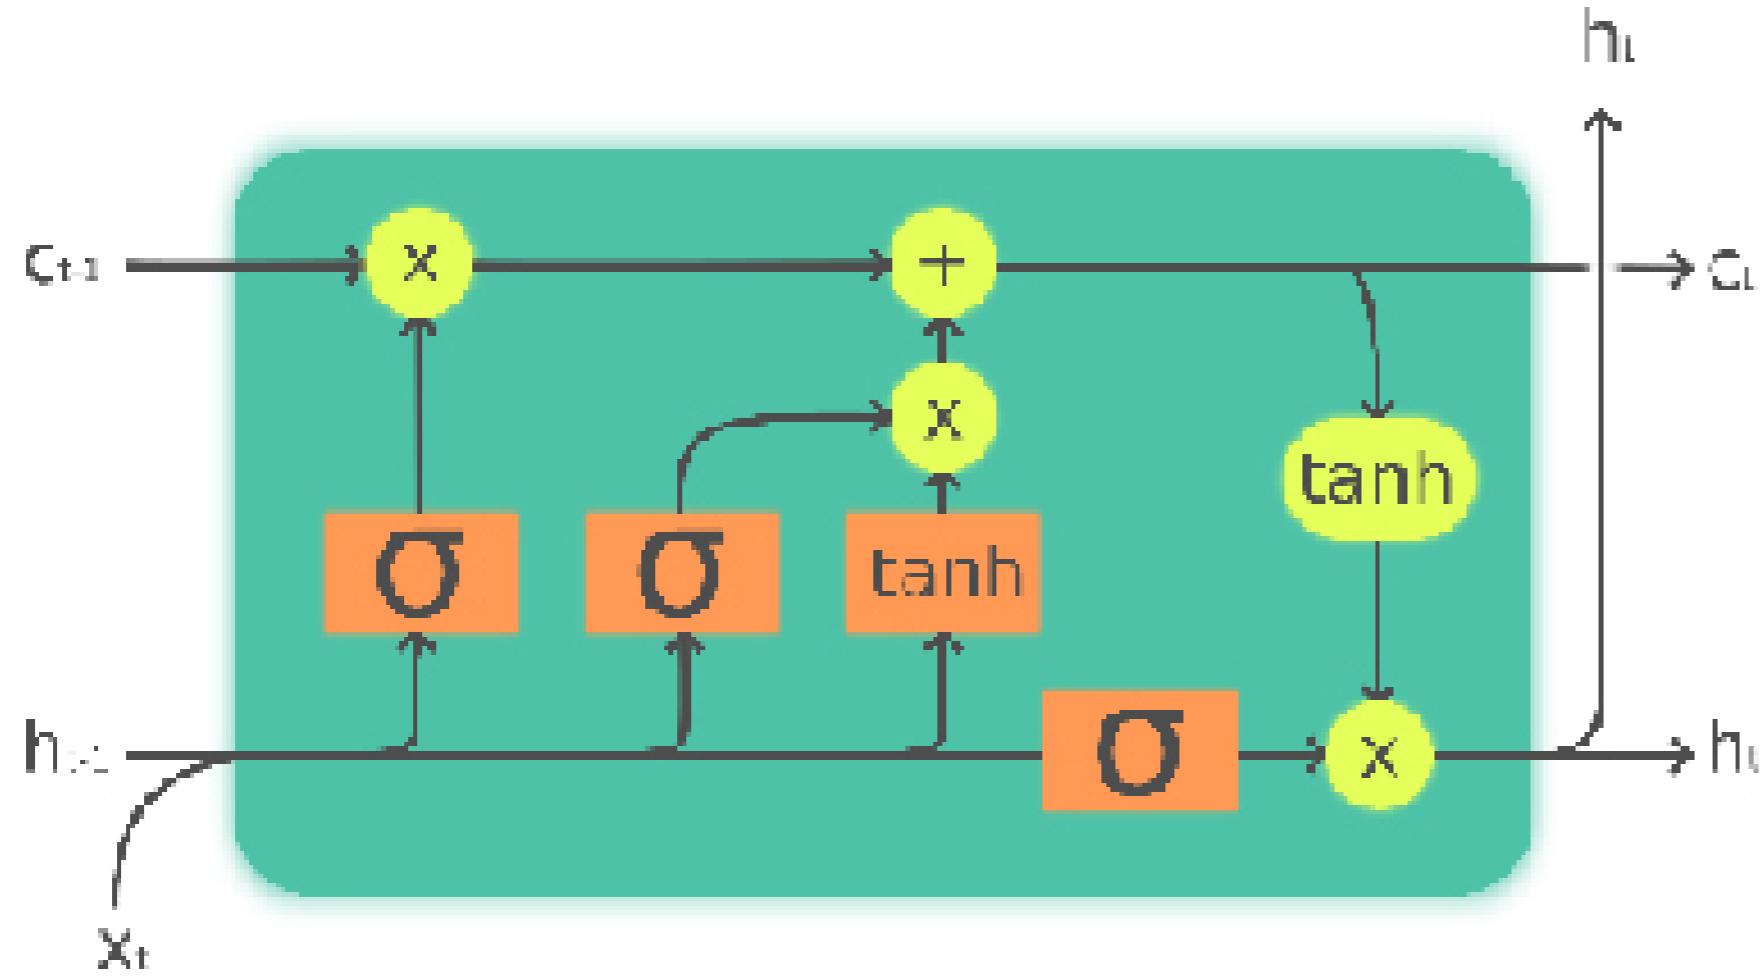


"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."

¹ Karpathy, A., & Fei-Fei, L. (2015). Deep visual-semantic alignments for generating image descriptions.

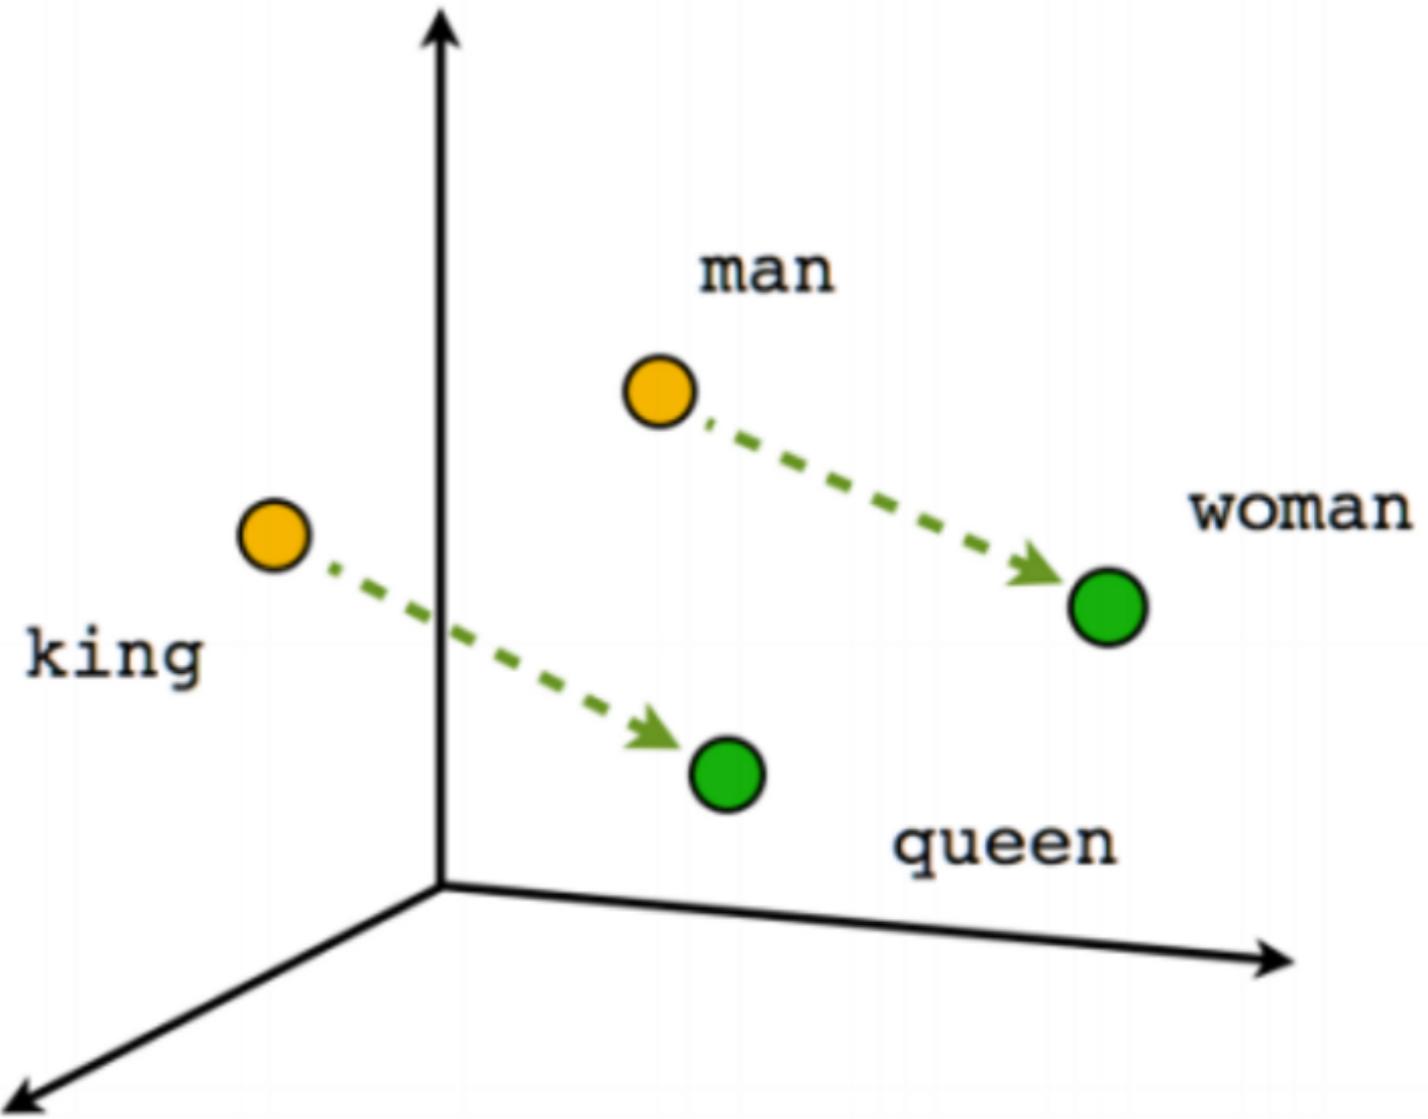
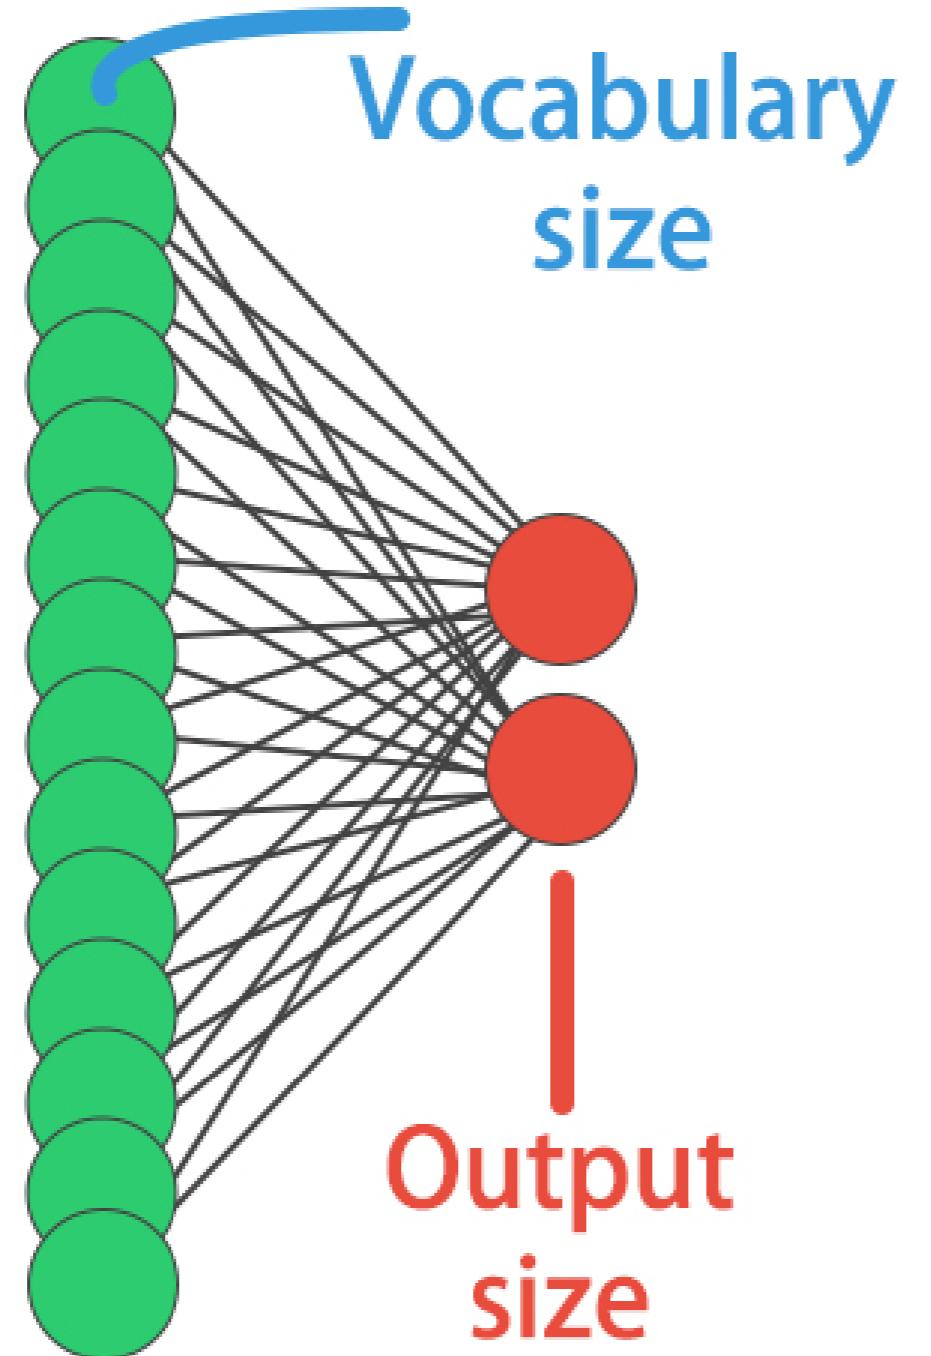


+ Text

this is a sentence



42 11 23 1



```
text = 'Hi this is a small sentence'

# We choose a sequence length
seq_len = 3

# Split text into a list of words
words = text.split()
```

```
['Hi', 'this', 'is', 'a', 'small', 'sentence']
```

```
# Make lines
lines = []
for i in range(seq_len, len(words) + 1):
    line = ' '.join(words[i-seq_len:i])
    lines.append(line)
```

```
['Hi this is', 'this is a', 'is a small', 'a small sentence']
```

```
# Import Tokenizer from keras preprocessing text
from keras.preprocessing.text import Tokenizer
# Instantiate Tokenizer
tokenizer = Tokenizer()
# Fit it on the previous lines
tokenizer.fit_on_texts(lines)
# Turn the lines into numeric sequences
sequences = tokenizer.texts_to_sequences(lines)
```

```
array([[5, 3, 1], [3, 1, 2], [1, 2, 4], [2, 4, 6]])
```

```
print(tokenizer.index_word)
```

```
{1: 'is', 2: 'a', 3: 'this', 4: 'small', 5: 'hi', 6: 'sentence'}
```

```
# Import Dense, LSTM and Embedding layers
from keras.layers import Dense, LSTM, Embedding
model = Sequential()
# Vocabulary size
vocab_size = len(tokenizer.index_word) + 1
# Starting with an embedding layer
model.add(Embedding(input_dim=vocab_size, output_dim=8, input_length=2))
# Adding an LSTM layer
model.add(LSTM(8))

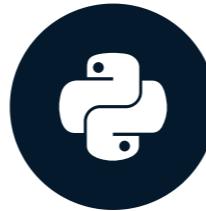
# Adding a Dense hidden layer
model.add(Dense(8, activation='relu'))
# Adding an output layer with softmax
model.add(Dense(vocab_size, activation='softmax'))
```

Let's do it!

INTRODUCTION TO DEEP LEARNING WITH KERAS

You're done!

INTRODUCTION TO DEEP LEARNING WITH KERAS



Miguel Esteban

Data Scientist & Founder

Congratulations!



You've learned a lot



What you've learned

- Basics of neural networks
- Building sequential models
- Building models for regression
- Approaching binary classification, multi-class and multi-label problems with neural networks
- Activation functions
- Hyperparameter optimization
- Autoencoders
- De-noising images
- CNN concepts
- Use pre-trained models
- Visualize convolutions
- LSTMs concepts
- Work with LSTMs and text
- All this by using many different datasets and learning a lot of Keras utility functions.

What could you learn next?

- Go deeper into CNNs
- Go deeper into LSTMs
- Keras Functional API
- Models that share layers, models with several branches
- GANs: Generative Adversarial Networks
- Deeplearning projects



It's sad to say goodbye
my friend

Have a good one!

INTRODUCTION TO DEEP LEARNING WITH KERAS